

## Category 1: Preprocessing & Data Strategy

**Q1: You used StandardScaler before training the KNN model. Why is scaling mandatory for KNN, whereas it might not be for a Decision Tree?**

- **Answer:** KNN calculates the mathematical **distance** (Euclidean distance) between points to classify them.
  - If one feature is measured in large numbers (e.g., Petal Length = 50mm) and another in small numbers (e.g., Sepal Width = 3mm), the distance formula will be biased toward the larger number.
  - Scaling puts all features on the same scale (Mean=0, Variance=1) so the model treats them equally. Decision Trees work on thresholds (rules), so they don't care about the scale.

**Q2: In your train\_test\_split, you used the argument stratify=y. What does this do, and why is it important?**

- **Answer:** Stratification ensures that the proportion of classes (Setosa, Versicolor, Virginica) in the **Training** set and the **Test** set is exactly the same as the original dataset.
  - Without this, random splitting might put all the "Setosas" in the training set and none in the test set, leading to a biased evaluation.

---

## Category 2: Principal Component Analysis (PCA)

**Q3: You used PCA to reduce the data to 2 dimensions and stated it captured "95.8% of the variance." What does that mean in plain English?**

- **Answer:** It means that by looking at just those 2 generated components (PC1 and PC2), we still possess **95.8% of the information** contained in the original 4 features. We threw away 2 dimensions but lost almost no useful data. This confirms that the original features were highly correlated (redundant).

**Q4: Looking at your code, did you use the PCA-transformed data for the actual KNN training, or just for the visualization?**

- **Answer:** (*Check your code carefully here!*)
  - I used PCA **only for visualization** (Step 3).
  - For the actual model training (Step 4 & 5), I went back to X\_raw (the original 4 features).
  - *Follow-up:* I could have used the PCA data for training to speed up the model, but since the dataset is small (150 rows), training on 4 features was fast enough, so I kept the original features to maximize accuracy.

---

## Category 3: The Model (KNN & GridSearchCV)

**Q5: You used GridSearchCV to find the best K. Why not just pick K=1 or K=50 manually?**

- **Answer:**

- **K=1 (Overfitting):** The model captures noise. If a generic flower has a weird mutation, the model memorizes it.
- **K=50 (Underfitting):** The model becomes too "blunt." It averages out too many neighbors and loses the ability to see distinct local clusters.
- GridSearchCV automates the testing of all these values (Cross-Validation) to find the mathematical "sweet spot" (which was K=5).

**Q6: Your final accuracy was 93.33%. Where do you think the 6.7% error came from?**

- **Answer:** Based on the **EDA (Pairplot)** and the **PCA Scatterplot**, the *Setosa* class is perfectly separated. However, *Versicolor* and *Virginica* have a slight physical overlap. The errors definitely occurred in that overlap region where the algorithm couldn't distinguish between a large *Versicolor* and a small *Virginica*.
- 

#### **Category 4: Troubleshooting & Code Proficiency**

**Q7: At the very end of your output, there are UserWarnings saying "X does not have valid feature names." Why did this happen?**

- **Answer:**

- The model was trained on a **Pandas DataFrame** (which has column names like 'sepal\_length', etc.).
- In my predict\_flower function, I passed a raw **NumPy array** (np.array([[...]])).
- The scaler/model complained because it expected column names but didn't get them.
- *Fix:* To fix this in production, I would convert the new input data into a Pandas DataFrame with the correct column names before passing it to .predict().

**Q8: If you deployed this model and new data came in with "Petal Length" in *Inches* instead of *Centimeters*, what would happen?**

- **Answer:** The model would break or give wrong predictions. Because the model relies on the StandardScaler fitted on the original CM data, inputting Inches (which are smaller numbers) would make the model think the flower is tiny. Data validation/unit conversion is required before the data hits the pipeline.